

FERNANDO LEONARDO CUEVAS

Senior Analytics Engineer

(718) 517-0680 | fernandoleonardocuevas@gmail.com | linkedin.com/in/fernandoleonardocuevas | github.com/fernandolcuevas

SUMMARY

Senior analytics engineer who turns raw data into the modeled, tested datasets and metrics a business runs on. Sole owner of a daily 10M+ record platform for dozens of Fortune 500 brands, and shipped production LLM segmentation on Gemini (Mosaic AI) benchmarked against a classical baseline. Five-plus years across Databricks, BigQuery, and AWS; B.S. Applied Mathematics & Cryptography.

EXPERIENCE

Senior Data Analyst - Traackr

Nov 2024 - Apr 2026

Remote, NYC

- Architected and solely owned the analytics orchestration platform on Databricks (API extraction, Unity Catalog transformations, enrichment, and delivery), with a self-service wizard that lets non-technical stakeholders run campaign, spend, and audience pipelines themselves - supporting \$5M+ in annual contract value.
- Standardized influencer and audience data models in Unity Catalog and Delta Lake with fact/dimension patterns, cutting recurring data-quality escalations.
- Built the campaign and spend metrics layer - multi-currency normalization rolled up to cost KPIs (CPP, CPE, CPVV, CPSV, CPME) consumed directly in client reporting.
- Migrated brand-and-product classification from 12 hours of manual work to a 2-minute distributed PySpark job - NLP feature extraction across millions of influencer records per cycle, validated against the senior analyst's golden-set output with full sample comparison.
- Built production LLM segmentation on Gemini via the Databricks Mosaic AI Foundation Model API - classifying ~65K posts a year across a two-year corpus and reaching 98% accuracy against a logistic-regression baseline - replacing the manual segmentation process, with batch handling, rate-limit management, and output validation.
- Built reconciliation checks that caught production data-loading gaps a shipped client deliverable had already run on - now a standard QA step on the workflow.
- Cut deliverable turnaround ~50% by re-architecting multi-stage ELT pipelines with rerunnable, failure-isolated execution, driving full-rerun-from-scratch incidents to near-zero across pipelines spanning tens of terabytes.
- Technical escalation point for a 5-person analytics team on the most complex, time-sensitive deliverables - including Samsung Unpacked launch tracking - and mentored analysts through new patterns, runbooks, and QA standards.

Data Scientist - Buxton (now Audiense)

May 2021 - Nov 2024

Remote, Ft. Worth, TX

- Powered "Personas," an LLM-and-embeddings system that generated client-specific demographic persona sets - ran the LLM execution and moved embeddings from Redshift to local compute for the final generation step, on AWS (S3, Redshift, Glue, Lambda).
- Reduced BigQuery cost and runtime ~15-20% on 150TB+ recurring operations using clustering and partitioning strategies.
- Owned a monthly foot-traffic export for a national retail client - ranking locations by month-over-month traffic across a portfolio of sites - automated end to end with Alteryx, GCP, MySQL, and SQL Server.
- Built ELT pipeline work on the team-owned, multi-source pipeline (Excel, Alteryx, Experian, BigQuery, GCS) behind a client-facing web app that generated \$1M+ in ARR.

SELECTED PROJECTS

- **LLM Influencer Segmenter** - two-stage segmentation pipeline pairing embedding-based retrieval with an LLM verification pass, with Unity Catalog integration and a full, typed pytest suite (59 passing tests). Built with production discipline - modular and reproducible - rather than as a notebook.
- **OCR Agent (OCR_GPT)** - folder-watching desktop agent that runs Tesseract OCR, queries an LLM, and surfaces the response; containerized with Docker and published to Docker Hub, with queue-based threading and typed error handling.

TECHNICAL SKILLS

Analytics Engineering: Dimensional (fact/dimension) modeling, modular SQL, ELT / ETL transformation pipelines, semantic / metrics layers, testing & reconciliation, self-service BI enablement, schema design

Languages: SQL, Python, PySpark, Spark SQL, T-SQL, Bash

Warehouses & Platforms: Databricks (Unity Catalog, Delta Lake, Jobs, Notebooks, Widgets, Mosaic AI), BigQuery, Redshift, SQL Server, MySQL

Data Engineering: End-to-end orchestration, rerunnable / failure-isolated execution, query optimization (clustering, partitioning), AWS (S3, Glue, Lambda), GCS

AI / ML: Production LLM integration (Gemini via Mosaic AI Foundation Model API), embeddings, prompt engineering, scikit-learn, output validation, cost/latency tradeoffs