

FERNANDO LEONARDO CUEVAS

Senior Analytics Engineer

(718) 517-0680 | fernandoleonardocuevas@gmail.com | linkedin.com/in/fernando-c-4325661b6 | github.com/fernandolcuevas

SUMMARY

Senior analytics engineer who turns raw data into the modeled, tested datasets and metrics a business runs on. Sole owner of a daily 10M+ record platform for dozens of Fortune 500 brands, and shipped production LLM segmentation on Gemini (Mosaic AI) benchmarked against a classical baseline. Five-plus years across Databricks, BigQuery, and AWS; B.S. Applied Mathematics and Cryptography.

EXPERIENCE

Senior Data Analyst, Traackr | Remote, NYC | Nov 2024 - Apr 2026

- Architected and solely owned the analytics orchestration platform on Databricks (API extraction, Unity Catalog transformations, enrichment, and delivery), with a self-service wizard that lets non-technical stakeholders run campaign, spend, and audience pipelines themselves, supporting \$5M+ in annual contract value.
- Standardized influencer and audience data models in Unity Catalog and Delta Lake with fact/dimension patterns, cutting recurring data-quality escalations.
- Built the campaign and spend metrics layer: multi-currency normalization rolled up to cost KPIs (CPP, CPE, CPVV, CPSV, CPME) consumed directly in client reporting.
- Migrated brand-and-product classification from 12 hours of manual work to a 2-minute distributed PySpark job, running NLP feature extraction across millions of influencer records per cycle, validated against the senior analyst's golden-set output with full sample comparison.
- Built production LLM segmentation on Gemini via the Databricks Mosaic AI Foundation Model API, classifying ~65K posts a year across a two-year corpus, against a logistic-regression baseline and a 98% accuracy target, replacing the manual segmentation process, with batch handling, rate-limit management, and output validation.
- Built reconciliation checks that caught production data-loading gaps a shipped client deliverable had already run on, now a standard QA step on the workflow.
- Cut deliverable turnaround ~50% by re-architecting multi-stage pipelines with rerunnable, failure-isolated execution, driving full-rerun-from-scratch incidents to near-zero across pipelines spanning tens of terabytes.
- Technical escalation point for a 5-person analytics team on the most complex, time-sensitive deliverables (including Samsung Unpacked launch tracking), and mentored analysts through new patterns, runbooks, and QA standards.

Data Scientist, Buxton (now Audiense) | Remote, Ft. Worth, TX | May 2021 - Nov 2024

- Powered "Personas," an LLM-and-embeddings system that generated client-specific demographic persona sets, running the LLM execution and moving embeddings from Redshift to local compute for the final generation step, on AWS (S3, Redshift, Glue, Lambda).
- Reduced BigQuery cost and runtime ~15-20% on 150TB+ recurring operations using clustering and partitioning strategies.
- Owned a monthly foot-traffic export for a national retail client, ranking locations by month-over-month traffic across a portfolio of sites, automated end to end with Alteryx, GCP, MySQL, and SQL Server.
- Built ELT pipeline work on the team-owned, multi-source pipeline (Excel, Alteryx, Experian, BigQuery, GCS) behind a client-facing web application that generated \$1M+ in ARR.

SELECTED PROJECT

LLM Influencer Segmenter: two-stage segmentation pipeline pairing embedding-based retrieval with an LLM verification pass, with Unity Catalog integration and a full, typed pytest suite. Built with production discipline (modular and reproducible) rather than as a notebook.

TECHNICAL SKILLS

Analytics Engineering: Dimensional (fact/dimension) modeling, modular SQL, transformation pipelines, semantic and metrics layers, testing and reconciliation, self-service BI enablement, schema design

Languages: SQL, Python, PySpark, Spark SQL, T-SQL, Bash

Warehouses and Platforms: Databricks (Unity Catalog, Delta Lake, Jobs, Notebooks, Widgets, Mosaic AI), BigQuery, Redshift, SQL Server, MySQL

Data Engineering: End-to-end orchestration, ETL/ELT, rerunnable and failure-isolated execution, query optimization (clustering, partitioning), AWS (S3, Glue, Lambda), GCS

AI and ML: Production LLM integration (Gemini via Mosaic AI Foundation Model API), embeddings, prompt engineering, classical ML baselines, output validation, cost and latency tradeoffs

DevOps: Git, GitHub Actions, CI/CD, Docker, Azure DevOps, shell scripting, operational runbooks

BI: Tableau, Alteryx

Libraries: Pandas, NumPy, Dask, Flask, scikit-learn, pytest, OpenAI SDK

EDUCATION AND CERTIFICATIONS

B.S. Applied Mathematics and Cryptography, John Jay College of Criminal Justice (CUNY), New York, NY | Dec 2020

Certified Alteryx Designer Core | Nov 2021