

FERNANDO LEONARDO CUEVAS

Senior Data Engineer

(718) 517-0680 • fernandoleonardocuevas@gmail.com • [linkedin.com/in/fernando-c-4325661b6](https://www.linkedin.com/in/fernando-c-4325661b6) • github.com/fernandolcuevas

SUMMARY

Senior data engineer and **sole owner of a daily 10M+ record production platform** on Databricks, supporting **\$5–15M ACV across 25–75 Fortune 500 brands**. Integrated **Gemini via the Mosaic AI Foundation Model API** into the production pipeline at 10M+ scale, and migrated a **12-hour manual classification process to a 2-minute distributed PySpark job**. 5+ years across Databricks, BigQuery, and AWS; B.S. Applied Mathematics & Cryptography.

EXPERIENCE

Senior Data Engineer / Analytics Engineer — Traackr

Nov 2024 – Apr 2026

Remote · NYC

- Architected an **end-to-end analytics orchestration platform** on Databricks — API extraction, Unity Catalog transformations, multi-stage enrichment, and Google Drive delivery — with a self-service wizard for non-technical stakeholders to run campaign, spend, and audience pipelines. **Sole owner end-to-end**.
- Owned daily production pipelines processing **10M+ records and 10–100TB** across **25–75 Fortune 500 clients**, supporting **\$5–15M in annual contract value**.
- Migrated brand-and-product classification from **12 hours of manual work to a 2-minute distributed PySpark job** — NLP feature extraction across 5–10M influencer records per cycle, validated against the senior analyst's golden-set output with full sample comparison.
- Integrated **Gemini via the Databricks Mosaic AI Foundation Model API** into the production audience pipeline at **10M+ scale** — prompt-driven enrichment, segmentation, and false-positive correction, with batch handling, rate-limit management, and output validation.
- Cut deliverable turnaround **~50%** by re-architecting multi-stage pipelines (refresh → export → hydration → packaging) with rerunnable, failure-isolated execution — cutting full-rerun-from-scratch incidents to near-zero.
- Standardized influencer and audience data models in **Unity Catalog and Delta Lake** with fact/dimension patterns, improving downstream correctness and reducing recurring data-quality escalations.
- Built **reconciliation checks** that caught production data-loading gaps a client-impacting deliverable had already shipped on — preventing further bad runs and becoming a standard QA step on the workflow.
- Served as **technical escalation point** for a 5-person analytics team on the most complex, time-sensitive deliverables (including Samsung Unpacked launch tracking), and mentored analysts through new patterns, runbooks, and tooling.

Analytics Developer / Analytical Engineer — Buxton (now Audiense)

May 2021 – Nov 2024

Remote · Ft. Worth, TX

- Reduced **BigQuery cost and runtime ~15–20%** on 150TB+ recurring operations using clustering and partitioning strategies.
- Built a multi-source production data pipeline (Excel, Alteryx, Experian, BigQuery, GCS) powering a client-facing web application that generated **\$1M+ in ARR**.
- Engineered **AWS-based ETL** with S3, Redshift, Glue, and Lambda — production data movement and transformation across multiple client accounts.

SELECTED PROJECT

OCR Agent (OCR_GPT) — desktop agent that watches a folder for screenshots, runs Tesseract OCR, queries an LLM, and surfaces the response. Queue-based threading bridge for safe UI updates; env-var configuration; typed error handling.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, Spark SQL, T-SQL, Bash

Data Platforms: Databricks (Unity Catalog, Delta Lake, Jobs, Notebooks, Widgets, Mosaic AI, distributed PySpark), BigQuery, AWS (S3, Redshift, Glue, Lambda), GCS, SQL Server

Data Engineering: End-to-end orchestration, multi-stage pipeline design, rerunnable / failure-isolated execution, fact/dimension modeling, query optimization (clustering, partitioning), schema design, modular SQL, reconciliation & QA

AI / LLM: Production LLM integration (Gemini via Mosaic AI Foundation Model API), prompt engineering at scale, batch handling, output validation, cost/latency tradeoffs, LLM agent patterns

DevOps: Git, GitHub Actions, CI/CD, Azure DevOps, shell scripting, operational runbooks

Libraries: Pandas, NumPy, Dask, Flask, OpenAI SDK

EDUCATION & CERTIFICATIONS

B.S. Applied Mathematics & Cryptography — John Jay College of Criminal Justice (CUNY), New York, NY Dec 2020

Certified Alteryx Designer Core

Nov 2021